

CS544


LESSON 8

MONGODB

SPRING MONGO

Today's requirements on databases

- Big data (large datasets)
- Agility
- Unstructured/ semi structured data




Database problems

- Too much data
 - The data does not fit anymore on one node




Database Scaling

■ Vertical scaling




● Horizontal scaling



Brewer's CAP Theorem

- A distributed system can support only two of the following characteristics



Consistency

- Strict consistency
 - The data that I read is always correct
 - You never loose data
- Eventual consistency
 - The data might not be correct
 - But will eventually become correct

Problems with relational databases

- Scaling writes are is very difficult and limited
 - Vertical scaling is limited and is expensive
 - Horizontal scaling is limited and is complex
 - Queries work only within shards
 - Strict consistency and partition tolerance leads to availability problems

A relational database is hard to scale

Problems with relational databases

- The schema in a database is fixed
- Schema evolution
 - Adding attributes to an object => have to add columns to table
 - You need to do a migration project
 - Application downtime ...

A relational database is hard to change

Problems with relational databases

- Relational schema doesn't easily handle unstructured and semi-structured data
 - Emails
 - Tweets
 - Pictures
 - Audio
 - Movies
 - Text

Unstructured data

The university has 5600 students.
John's ID is number 1, he is 18 years old and already holds a B.Sc. degree.
David's ID is number 2, he is 31 years old and holds a Ph.D. degree. Robert's ID is number 3, he is 51 years old and also holds the same degree as David, a Ph.D. degree.

Semi-structured data

```
<University>
  <Student ID="1">
    <Name>John</Name>
    <Age>18</Age>
    <Degree>B.Sc.</Degree>
  </Student>
  <Student ID="2">
    <Name>David</Name>
    <Age>31</Age>
    <Degree>Ph.D. </Degree>
  </Student>
  ....
</University>
```

Structured data


ID	Name	Age	Degree
1	John	18	B.Sc.
2	David	31	Ph.D.
3	Robert	51	Ph.D.
4	Rick	26	M.Sc.
5	Michael	19	B.Sc.

A relational database does not handle unstructured and semi structured data very well

NoSQL characteristics


- Key-value store
- No fixed schema
- Can scale (almost) unlimited
 - Eventual consistency

Different types of databases



MongoDB

- Document database
- Fast
- Can handle large datasets



MongoDB

RDBMS		MongoDB
Database	→	Database
Table	→	Collection
Row	→	Document
Index	→	Index
Join	→	Embedded Document
Foreign Key	→	Reference



Document data model (JSON)

Relational - Tables

Customer ID	First Name	Last Name	City
0	John	Doe	New York
1	Mark	Smith	San Francisco
2	Jay	Black	Newark
3	Meagan	White	London
4	Edward	Daniels	Boston

Account Number	Branch ID	Account Type	Customer ID
10	100	Checking	0
11	101	Savings	0
12	101	IRA	0
13	200	Checking	1
14	200	Savings	1
15	201	IRA	2



Document - Collections

```
{  customer_id : 1,
  first_name  : "Mark",
  last_name   : "Smith",
  city        : "San Francisco",
  accounts   : [ {
    account_number : 13,
    branch_ID      : 200,
    account_type    : "Checking"
  },
  {
    account_number : 14,
    branch_ID      : 200,
    account_type    : "IRA",
    beneficiaries  : [...]
  } ]
}
```

Documents are rich structures

```
{  
  category: "glove",  
  model: "PRO112PT",  
  name: "Air Elite",  
  brand: "Rawlings",  
  price: 229.99,  
  available: Date("2013-03-31"),  
  position: ["infield", "outfield", "pitcher"]  
}
```

Fields can contain arrays

Documents are rich structures

```
{  
  category: "glove",  
  model: "PRO112PT",  
  name: "Air Elite",  
  brand: "Rawlings",  
  price: 229.99,  
  available: Date("2013-03-31"),  
  position: ["infield", "outfield", "pitcher"],  
  endorsed: {name: "Ryan Howard",  
             team: "Phillies",  
             position: "first base"},  
}
```

} Fields can contain
sub-documents

Documents are rich structures

```
{  
  category: "glove",  
  model: "PRO112PT",  
  name: "Air Elite",  
  brand: "Rawlings",  
  price: 229.99,  
  available: Date("2013-03-31"),  
  position: ["infield", "outfield", "pitcher"],  
  endorsed: {name: "Ryan Howard",  
             team: "Phillies",  
             position: "first base"},  
  history: [{date: Date("2013-03-31"), price: 279.99},  
            {date: Date("2013-06-01"), price: 259.79},  
            {date: Date("2013-08-15"), price: 229.99}]  
}
```

Fields can contain
an array of sub-
documents

Documents are flexible

{

category: bat,
model: B1403E,
name: Air Elite,
brand: "Rip-IT",
price: 399.99

diameter: "2 5/8",
barrel: R2 Alloy,
handle: R2

}

{

category: glove,
model: PRO112PT,
name: Air Elite,
brand: "Rawlings",
price: "229.99"

size: 11.25,
position: outfield,
pattern: "Pro taper",
material: leather,
color: black

}

BSON

```
{ author: 'joe',  
  created : new Date('03/28/2009'),  
  title : 'Yet another blog post',  
  text : 'Here is the text...',  
  tags : [ 'example', 'joe' ],  
  comments : [  
    { author: 'jim',  
      comment: 'I disagree'  
    },  
    { author: 'nancy',  
      comment: 'Good post'  
    }  
  ]  
}
```



Remember it is stored in binary formats (BSON)

```
"\x16\x00\x00\x00\x02hello\x00  
\x06\x00\x00\x00world\x00\x00"  
  
"1\x00\x00\x00\x04BSON\x00&\x00  
\x00\x00\x020\x00\x08\x00\x00  
\x00awesome\x00\x011\x00333333  
\x14@\x102\x00\xc2\x07\x00\x00  
\x00\x00"
```

Find() method

SQL SELECT Statements

SELECT * **FROM** users

SELECT id, user_id, status **FROM** users

SELECT user_id, status **FROM** users

SELECT * **FROM** users **WHERE** status = "A"

SELECT user_id, status **FROM** users **WHERE** status = "A"

SELECT * **FROM** users **WHERE** status != "A"

SELECT * **FROM** users **WHERE** status = "A" **AND** age = 50

SELECT * **FROM** users **WHERE** status = "A" **OR** age = 50

SELECT * **FROM** users **WHERE** age > 25

MongoDB find() Statements

db.users.find()

db.users.find({}, { user_id: 1, status: 1 })

db.users.find({}, { user_id: 1, status: 1, _id: 0 })

db.users.find({ status: "A" })

db.users.find({ status: "A" }, { user_id: 1, status: 1, _id: 0 })

db.users.find({ status: { \$ne: "A" } })

db.users.find({ status: "A", age: 50 })

db.users.find({ \$or: [{ status: "A" }, { age: 50 }] })

db.users.find({ age: { \$gt: 25 } })

Find() method



```
SELECT * FROM users WHERE age < 25
```

```
db.users.find( { age: { $lt: 25 } } )
```

```
SELECT * FROM users WHERE age > 25 AND age <= 50
```

```
db.users.find( { age: { $gt: 25, $lte: 50 } } )
```

```
SELECT * FROM users WHERE user_id like "%bc%"
```

```
db.users.find( { user_id: /bc/ } )
```

```
SELECT * FROM users WHERE user_id like "bc%"
```

```
db.users.find( { user_id: /^bc/ } )
```

```
SELECT * FROM users WHERE status = "A" ORDER BY user_id ASC
```

```
db.users.find( { status: "A" } ).sort( { user_id: 1 } )
```

```
SELECT * FROM users WHERE status = "A" ORDER BY user_id DESC
```

```
db.users.find( { status: "A" } ).sort( { user_id: -1 } )
```


```
SELECT COUNT(*) FROM users
```

```
db.users.count()  
or  
db.users.find().count()
```

```
SELECT COUNT(user_id) FROM users
```

```
db.users.count( { user_id: { $exists: true } } )  
or  
db.users.find( { user_id: { $exists: true } } ).count()
```

Schema free



The image displays five JSON documents in a light blue, torn-paper style font on a dark background. The documents are arranged in three columns: the first column has one document, the second has two, and the third has two. Each document represents a different person with varying fields, illustrating a schema-free database structure.

```
{name: "will",  
  eyes: "blue",  
  birthplace: "NY",  
  aliases: ["bill", "la  
ciacco"],  
  gender: "???",  
  boss: "ben"}
```

```
{name: "jeff",  
  eyes: "blue",  
  height: 72,  
  boss: "ben"}
```

```
{name: "brendan",  
  aliases: ["el diablo"]}
```

```
{name: "ben",  
  hat: "yes"}
```

```
{name: "matt",  
  pizza: "DiGiorno",  
  height: 72,  
  boss: 555.555.1212}
```

Main point

- MongoDB is a document database that stores whole documents (including embedded data) in a collection. This gives data redundancy, but makes the data access very fast.

Science of Consciousness: The Unified Field is the source of all relative creation where there is no redundancy or loss of performance.


Spring Mongo libraries

```
<dependency>  
  <groupId>org.springframework.boot</groupId>  
  <artifactId>spring-boot-starter-data-mongodb</artifactId>  
</dependency>
```

The Mongo Documents

```
@Document
public class Customer {
    @Id
    private int customerNumber;
    private String name;
    private String email;
    private String phone;
    private CreditCard creditCard;
```

```
public class CreditCard {
    private String cardNumber;
    private String type;
    private String validationDate;
```



The repository

@Repository

```
public interface CustomerRepository extends MongoRepository<Customer, Integer> {  
    Customer findByPhone(String phone);  
    Customer findByEmail(String email);  
    List<Customer> findByCreditCardType(String type);  
  
    @Query("{email : #{#email}}")  
    Customer findCustomerWithEmail(@Param("email") String email);  
}
```

application.properties

```
spring.data.mongodb.host=localhost  
spring.data.mongodb.port=27017  
spring.data.mongodb.database=testdb
```

The application(1/2)

```
public class Application implements CommandLineRunner {
```

```
    @Autowired
```

```
    private CustomerRepository customerRepository;
```

```
    public static void main(String[] args) {
```

```
        SpringApplication.run(Application.class, args);
```

```
    }
```

```
    @Override
```

```
    public void run(String... args) throws Exception {
```

```
        // create customer
```

```
        Customer customer = new Customer(101, "John doe", "johnd@acme.com", "0622341678");
```

```
        CreditCard creditCard = new CreditCard("12324564321", "Visa", "11/23");
```

```
        customer.setCreditCard(creditCard);
```

```
        customerRepository.save(customer);
```

```
        customer = new Customer(109, "John Jones", "jones@acme.com", "0624321234");
```

```
        creditCard = new CreditCard("657483342", "Visa", "09/23");
```

```
        customer.setCreditCard(creditCard);
```

```
        customerRepository.save(customer);
```

```
        customer = new Customer(66, "James Johnson", "jj123@acme.com", "068633452");
```

```
        creditCard = new CreditCard("99876549876", "MasterCard", "01/24");
```

```
        customer.setCreditCard(creditCard);
```

```
        customerRepository.save(customer);
```

The application(2/2)

```
//get customers
System.out.println(customerRepository.findById(66).get());
System.out.println(customerRepository.findById(101).get());
System.out.println("-----All customers -----");
System.out.println(customerRepository.findAll());
//update customer
customer = customerRepository.findById(101).get();
customer.setEmail("jd@gmail.com");
customerRepository.save(customer);
System.out.println("-----find by phone -----");
System.out.println(customerRepository.findByPhone("0622341678"));
System.out.println("-----find by email -----");
System.out.println(customerRepository.findCustomerWithEmail("jj123@acme.com"));
System.out.println("-----find customers with a certain type of creditcard -----");
List<Customer> customers = customerRepository.findByCreditCardType("Visa");
for (Customer cust : customers){
    System.out.println(cust);
}
}
```

Embedded

testdb.customer

Documents

Aggregations

FILTER

{ field: 'value' }

ADD DATA

VIEW

>

_id: 101

name: "John doe"

email: "jd@gmail.com"

phone: "0622341678"

creditCard: Object

cardNumber: "12324564321"

type: "Visa"

validationDate: "11/23"

_class: "embedded.domain.Customer"

_id: 109

name: "John Jones"

email: "jones@acme.com"

phone: "0624321234"

creditCard: Object

cardNumber: "657483342"

type: "Visa"

validationDate: "09/23"

_class: "embedded.domain.Customer"

_id: 66

name: "James Johnson"

email: "jj123@acme.com"

phone: "068633452"

creditCard: Object

cardNumber: "99876549876"

type: "MasterCard"


validationDate: "01/24"

_class: "embedded.domain.Customer"

The Mongo Documents

```
@Document
public class Customer {
    @Id
    private int customerNumber;
    private String name;
    private String email;
    private String phone;
    @DBRef
    private CreditCard creditCard;
```

```
public class CreditCard {
    private String cardNumber;
    private String type;
    private String validationDate;
```



Reference

testdb.customer

Documents

Aggregations

Schema

FILTER

{ field: 'value' }

ADD DATA

VIEW

>

_id: 101

name: "John doe"

email: "jd@gmail.com"

phone: "0622341678"

creditCard: DBRef(creditCard, 12324564321, undefined)

_class: "reference.domain.Customer"

_id: 109

name: "John Jones"

email: "jones@acme.com"

phone: "0624321234"

creditCard: DBRef(creditCard, 657483342, undefined)

_class: "reference.domain.Customer"

_id: 66

name: "James Johnson"

email: "jj123@acme.com"

phone: "068633452"

creditCard: DBRef(creditCard, 99876549876, undefined)

_class: "reference.domain.Customer"

testdb.creditCard

Documents

Aggregations

FILTER

{ field: 'value' }

ADD DATA

VIEW

_id: "12324564321"

type: "Visa"

validationDate: "11/23"

_class: "reference.domain.CreditCard"

_id: "657483342"

type: "Visa"

validationDate: "09/23"

_class: "reference.domain.CreditCard"

_id: "99876549876"

type: "MasterCard"

validationDate: "01/24"

_class: "reference.domain.CreditCard"

One to many relations

```
public class Customer {  
    @Id  
    private int customerNumber;  
    private String name;  
    private String email;  
    private String phone;  
    private List<CreditCard> creditCards = new ArrayList<CreditCard>();  
}
```

```
public class CreditCard {  
    private String cardNumber;  
    private String type;  
    private String validationDate;  
}
```

The repository

@Repository

```
public interface CustomerRepository extends MongoRepository<Customer, Integer> {  
    Customer findByPhone(String phone);  
    List<Customer> findByCreditCardsType(String type);  
}
```

Connecting the parts of knowledge with the wholeness of knowledge

1. MongoDB is a document database where we store documents instead of relational data
2. Spring Boot Mongo makes it very easy to use the MongoDB in your application

-
3. **Transcendental consciousness** is the field where all intelligence resides.
 4. **Wholeness moving within itself:** In unity consciousness, one experiences that everything is an expression of one's own Self.

